

Tim auf der Landwehr
t.auf.der.landwehr@student.rug.nl
S2548682
Last revision: December 8, 2013

Learning from Data

Assignment 4

Summary

This is the assignment 4 in the course Learning from Data at the University of Groningen.

Exercise 1: arff-file generation for reuters-allcats.csv

The arff file generation script is provided in `reuters_cluster_arff_gen.py`. It is supposed to be used in the following manner. If you run in the first time, the parameter `--features` needs to be used to generate the features file.

```
Reuters-ARFF-Generator v1.0 by Tim auf der Landwehr

Usage: reuters_cluster_arff_gen.py
  Loads features from [reuters.features.json] and dialectology from
  [reuters-allcats.csv]
  and writes arff output to [reuters-allcats.arff]

--features
  Generates features to [reuters.features.json].
--in file
  Change name of the annotated input file.
--out file
  Change name of the arff output file.
```

Exercise 2: Running K-means

Running K-means on the arff file generated in exercise 1 by the command

```
java -cp weka.jar weka.clusterers.SimpleKMeans -c 1 -N 7 -t reuters-allcats.arff
```

leads to the following results

Clustered Instances

0	3 (0%)
1	1 (0%)
2	1 (0%)
3	465 (11%)
4	2 (0%)
5	2356 (58%)
6	1251 (31%)

Class attribute: @@class@@

Classes to Clusters:

0	1	2	3	4	5	6	<-- assigned to cluster
3	1	1	349	2	1293	280	Neg-
0	0	0	70	0	509	211	Pos-acq
0	0	0	5	0	29	1	Pos-coffee
0	0	0	39	0	490	751	Pos-earn
0	0	0	2	0	29	3	Pos-gold
0	0	0	0	0	0	4	Pos-heat
0	0	0	0	0	6	1	Pos-housing

```
Cluster 0 <-- No class
Cluster 1 <-- No class
Cluster 2 <-- No class
Cluster 3 <-- Pos-acq
Cluster 4 <-- No class
Cluster 5 <-- Neg-
Cluster 6 <-- Pos-earn
```

Incorrectly clustered instances : 1965.0 48.1736 %

This result is actually not very satisfactory, but it was the best I could reach.

Exercise 3: Lexical classification of dutch dialects

I solved this task using all unicode letters from \u0020 to \u03E8. Therefore the features are totally independent from the input. This range was defined by me with the help of the script `check_unicode_in_dialectology.py`, which checks all the characters in an given file against a certain range of unicode characters and tries to find those that do not occur in the range. Using this on several of the dialectology files, I figured that the defined range is sufficient to match all characters that are used in these files.

To be able to judge the accuracy of this features, I annotated the file `035.txt` using the script `annotate_dialectology.py`, that simply adds the classes 'devil', 'lucifer' and 'no_class' to the each line depending on whether it starts with *d*, *l* or something else. This divides the data quite well into the two classes and sets 'no_class' for those line which do not contain any data (but '|').

This annotated file can be processed with the script `dialectology_arff_gen.py`, which generates an WEKA-conform arff file from the annotated input file.

The arff-file can be processed in weka using

```
java -cp weka-3-6-10/weka.jar weka.clusterers.SimpleKMeans -c 1 -N 3 -t
assignment4/035.arff
```

That leads to the following output

Clustered Instances

```
0      21 ( 6%)
1     231 ( 64%)
2     110 ( 30%)
```

Class attribute: @@class@@

Classes to Clusters:

```
0  1  2 <-- assigned to cluster
1 225 0 | devil
20  0 109 | lucifer
0   6  1 | no_class
```

Cluster 0 <-- No class

Cluster 1 <-- devil

```
Cluster 2 <-- lucifer
```

```
Incorrectly clustered instances : 28.0 7.7348 %
```

This is a quite good result, but I should mention, that a simple classifier that only uses the first letter of the line would work better in this case. But this approach will do better clustering other words.